

1. My proposed row key is hash(VIN) + “#” + ModelYear + “#” + Make + “#” + City. If the row key begins with an attribute like model year, new data will constantly be written to the same tablet, which will cause hotspotting. This is when one node gets overloaded. By hashing the VIN as the prefix of the row key, writes are distributed evenly across tablets, no node gets too much write traffic, and parallel writes work fine. The VIN is ideal because it is unique for each car.
2. Column families should group attributes that are accessed together. Bigtable stores column families separately which means this would improve the efficiency